

Meeting AI Report: Exeter.mp4

Date: 2026-07-29 17:04:40

Duration: 01:23:20

Total Segments: 10

We are to rewrite the given transcript, preserving the timestamps and format. The provided transcript has speaker turns with speaker IDs? Actually, the example shows "Speaker A (00:00:50 - 00:00:51)" etc. But in the given text, it's listed as [00:39:27.560 --> 00:39:29.240] for the first speaker, then "[00:39:32.000 --> 00:39:33.520]" etc. However, looking at the user's message, it seems the transcript is given as a series of lines with timestamps and speaker comments? But the instruction says "preserve the timestamps and format each speaker's turn exactly like this": [Text spoken by Speaker A]. So we need to extract the speaker turns and present them in that format.

But note: The user provided a transcript (from a video) where there are speakers talking, with timestamps. We need to format each speaker's turn as "Speaker X (timestamp1 - timestamp2)" followed by the spoken words. However, looking at the given data, it seems the speaker turns are listed without explicit speaker labels? Actually, in the example, we see "[00:39:27.560 --> 00:39:29.240]" which is likely a timestamp for the turn of a speaker. Then after that, there's text spoken by that speaker. But in the problem statement, it says "You MUST preserve the timestamps and format each speaker's turn exactly like this: [Text spoken by Speaker A]". So we need to parse the given transcript and convert it into the required format, where each turn is labeled with a speaker identifier (like Speaker 1, etc.) but in the given data, the speakers are referred to implicitly? Actually, looking at the raw input: It's a list of lines with timestamps and then text. The first line says "if you're able to come for an in-campus visit" and then next timestamp, and then "[00:39:27.560] ...". So it seems that the speaker is implied by context, but we should treat each timestamp as the turn of a speaker, and then capture what they say.

Given the provided transcript, we can assume that each line after the timestamp is spoken by a different speaker. The format should be: [speaker identifier] (timestamp) -> [spoken words]. But in the given data, the first speaker's turn is "[00:39:27.560] ...". So likely the "Speaker A" is that part after the timestamp, but we need to assume that each line is a turn of a speaker. However, the problem says "format each speaker's turn exactly like this: [Text spoken by Speaker A]", meaning the text should be the words spoken by that speaker, and the surrounding maybe the timestamp? Actually, looking at the example format in the user's instruction: it shows "[Text spoken by Speaker A]" which suggests the speaker is labeled, and then the actual text. But in the given transcript, after the timestamp there is the speech. So we should present each turn as [Speaker] (timestamp) followed by the spoken words? But in the example, for the first line, it's "[00:39:27.560] if you're able to come for an in-campus visit". That seems like the timestamp is part of the label and then the speech. However, typical transcript format would have speaker ID and then time interval and then the speaker's utterance. But here, they've combined timestamp with the label. So we can just follow: start with [Speaker], then timestamp, then the spoken words? But in the given data, after the timestamp there is a colon? Actually, looking at the raw input:

[00:39:27.560] if you're able to come for an in-campus visit

[00:39:29.240] and really see our students in action,
[00:39:32.000] ask them your questions, get in conversations with them,
[00:39:33.520] they would be delighted,
[00:39:38.080] and they are just such wonderful resources to each other.
[00:39:45.880] So if you have the opportunity to meet our students out there,
[00:39:49.320] I encourage you to do so.

- Well, thank you both.

And just to cap that off, it's okay to be nervous.

I get a lot of, a lot of, I have a lot of conversations
with families who start off, well, I'm a little nervous.

And I always say, it's okay, you should be nervous.

I think Eileen said it.

This is the most important thing in your life.

So it's okay to be nervous.

I think it's important.

And I'll quote one of my mentors,
that's the beauty of these schools, to make you nervous.

That's where the growth happens.

Be nervous a little bit.

And yeah, your student will do well.

Let them learn about the school,
let them find what excites them.

And yeah, yeah, let them drive the process.

And we appreciate you as parents
of you being so invested and taking the time
to be with us this evening, to ask your thoughtful questions.

Your students have asked us
a lot of thoughtful questions as well.

So if you've had a question
and your student has attended one of our in-person sessions,

don't be afraid to ask that student.

It might have been repeated already, but also ask our team.

We're here to help support you guys through this process.

But again, thank you all so much for being with us tonight.

We hope you enjoyed the session.

Again, this recording will be available

probably tomorrow on the student's applicant portal.

And thank you for being with us.

- Have a wonderful night.

Reach out to us if you need anything.

Thank you so much.

- Bye everyone, have a good night.

Then (upbeat music) ... etc.

So each of these is a turn spoken by someone, presumably the speaker's speech. But in terms of labeling speakers, we can assume that the first line after timestamp is Speaker A, then the next is Speaker B, etc., but here there are multiple speakers: the first speaker (maybe about in-campus visit), then another speaker (students or something), then a person saying "Well, thank you both." and then "And just to cap that off...", etc. But note: there's a line "- Well, thank you both." which might be from someone else, then "And just to cap that off..." is from the same? Actually, it's part of the flow. However, for formatting, we should present each speaker turn with their label. The instruction says "format each speaker's turn exactly like this: [Text spoken by Speaker A]", meaning we need to indicate who is speaking (speaker A), then timestamp and then the text.

But in the given data, the timestamps are separated by spaces and there are line breaks. So we can treat each of these as a separate turn of a speaker. We should assign labels: perhaps the first speaker's turn is "Speaker 1", then its content, then next speaker's turn "Speaker 2", etc. However, in the provided transcript, the timestamps are consecutive and the text after timestamp is the spoken words. So we can just present each as: [Speaker] (timestamp) -> [spoken words]. But note: the given format for the sample has "[00:39:27.560] if you're able to come for an in-campus visit". The "if..." seems like the speaker's statement, so we can extract that as the spoken word, and then the timestamp is the turn. So we need to preserve the timestamps and format each line as [Speaker] (timestamp) followed by the text? But the example shows "[00:39:27.560] if you're able to come for an in-campus visit". The "if you're able to come for an in-campus visit" is likely what Speaker A is saying, so we can consider that part as speaker A's turn. So we should list each of those as a separate entry with the timestamp and then the text.

Given that, we need to reconstruct the transcript by taking each segment between timestamps (the speaker turns) and extract the spoken words. However, note there are pauses or line breaks? The input is a sequence of lines with timestamps and possibly punctuation. We should parse them into turns: each turn is from timestamp1 to timestamp2, and in that turn, what is said? Actually, the given text has no newline after the timestamp? Wait, looking at the raw data:

[00:39:27.560] if you're able to come for an in-campus visit

[00:39:29.240] and really see our students in action,

[00:39:32.000] ask them your questions, get in conversations with them,

[00:39:33.520] they would be delighted,

[00:39:38.